

V2 recall-QPS Pareto: baselines (16-thread real) + ML router (per-query × 16)

